

[ceu-economics-and-business.github.io](https://github.io/ceu-economics-and-business)

# Data warehouse architectures. – ECBS 5146 SQL and Different Shapes of Data

7–8 minutes

---

## Overview

**Teaching:** 90 min

## Questions

- What is a Data Warehouse?
- As analyst, how can you create a simple analytical platform using SQL DB?

## Objectives

- Understanding data warehouse architectures
- Building a denormalized analytical data store
- Building an ETL pipeline using MySQL Triggers and Events
- Building data marts with MySQL View

## Keywords

#DATA WAREHOUSE ARCHITECTURE

#VIEWS

#TRIGGERS

#ETL

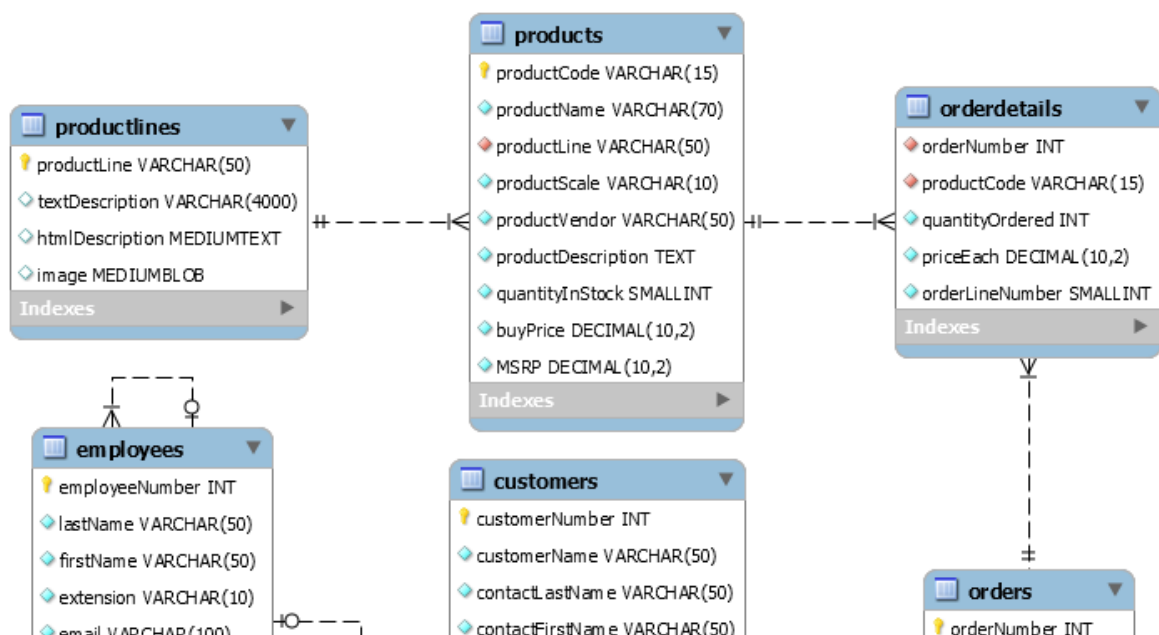
#DATA MARTS

## Table of Content

[Lecture PPTX](#)[Session setup](#)[Creating the analytical data store](#)[Events to schedule ETL jobs](#)[Trigger as ETL](#)[Data marts with Views](#)[Security with Views](#)[Term project](#)

## Session setup

No need to load new data, in this chapter we will use the same sample db we used in the last chapter:





## Creating the analytical data store

```
CREATE TABLE new_order LIKE orders;
```

```
DROP TABLE new_order;
```

```
CREATE TABLE new_order AS SELECT * FROM orders;
```

### Exercise1

Create a physical copy of birdstrikes with records where state is Oklahoma

We will use a query created in Homework 3. This creates a denormalized snapshot of the operational tables for product\_sales subject. We will embed the creation in a stored procedure.

```
DROP PROCEDURE IF EXISTS CreateProductSalesStore;
```

```
DELIMITER //
```

```
CREATE PROCEDURE CreateProductSalesStore()
BEGIN

    DROP TABLE IF EXISTS product_sales;

    CREATE TABLE product_sales AS
    SELECT
        orders.orderNumber AS SalesId,
        orderdetails.priceEach AS Price,
        orderdetails.quantityOrdered AS Unit,
        products.productName AS Product,
        products.productLine As Brand,
        customers.city As City,
        customers.country As Country,
        orders.orderDate AS Date,
        WEEK(orders.orderDate) as WeekOfYear
    FROM
        orders
    INNER JOIN
        orderdetails USING (orderNumber)
    INNER JOIN
        products USING (productCode)
    INNER JOIN
        customers USING (customerNumber)
    ORDER BY
        orderNumber,
        orderLineNumber;

END //
```

```
DELIMITER ;
```

```
CALL CreateProductSalesStore();
```

## Events to schedule ETL jobs

Event engine runs scheduled jobs/tasks. We can use it for scheduling ETL processes.

Basics on how to check the state of the scheduler. Check if scheduler is running

```
SHOW VARIABLES LIKE "event_scheduler";
```

Turn it on if not

```
SET GLOBAL event_scheduler = ON;
```

This is how you turn it OFF

```
SET GLOBAL event_scheduler = OFF;
```

Format:

```
CREATE EVENT [IF NOT EXISTS] event_name  
ON SCHEDULE schedule  
DO  
event_body
```

### Exercise 3

create a scheduler which writes the current time in messages in every second

Event which is calling CreateProductSalesStore every 1 minute in the next 1 hour.

```
DELIMITER $$

CREATE EVENT CreateProductSalesStoreEvent
ON SCHEDULE EVERY 1 MINUTE
STARTS CURRENT_TIMESTAMP
ENDS CURRENT_TIMESTAMP + INTERVAL 1 HOUR
DO
    BEGIN
        INSERT INTO messages SELECT
CONCAT('event:',NOW());
        CALL CreateProductSalesStore();
    END$$
DELIMITER ;
```

Listing all events stored in the schema

Deleting an event

```
DROP EVENT IF EXISTS CreateProductSalesStoreEvent;
```

## Trigger as ETL

Format:

```
DELIMITER $$

CREATE TRIGGER trigger_namex
    AFTER INSERT ON table_namex FOR EACH ROW
BEGIN
    -- statements
    -- NEW.orderNumber, NEW.productCode etc
END$$
```

```
DELIMITER ;
```

### Exercise3

Copy the birdstrikes structure into a new table called birdstrikes2.

Insert into birdstrikes2 the line where id is 10. Hints:

- Use the samples from Chapter2 for copy
- For insert use the format like: INSERT INTO bla SELECT blabla

Empty log table:

### The trigger

Creating a trigger which is activated if an insert is executed into orderdetails table. Once triggered will insert a new line in our previously created data store.

```
DROP TRIGGER IF EXISTS after_order_insert;

DELIMITER $$

CREATE TRIGGER after_order_insert
AFTER INSERT
ON orderdetails FOR EACH ROW
BEGIN

    -- log the order number of the newley
inserted order
    INSERT INTO messages SELECT CONCAT('new
orderNumber: ', NEW.orderNumber);
```

```
-- archive the order and associated table
entries to product_sales
INSERT INTO product_sales
SELECT
    orders.orderNumber AS SalesId,
    orderdetails.priceEach AS Price,
    orderdetails.quantityOrdered AS Unit,
    products.productName AS Product,
    products.productLine As Brand,
    customers.city As City,
    customers.country As Country,
    orders.orderDate AS Date,
    WEEK(orders.orderDate) as WeekOfYear
FROM
    orders
INNER JOIN
    orderdetails USING (orderNumber)
INNER JOIN
    products USING (productCode)
INNER JOIN
    customers USING (customerNumber)
WHERE orderNumber = NEW.orderNumber
ORDER BY
    orderNumber,
    orderLineNumber;

END $$

DELIMITER ;
```



E - Extract: Joining the tables for the operational layer is an extract operation

T - Transform: We don't have glamorous transformations here, only a WeekOfYear covering this part. Nevertheless, please note that you call a store procedure from trigger or even use procedural language to do transformation in the trigger itself.

L - Load: Inserting into product\_sales represents the load part of the ETL

### Activating the trigger

Listing the current state of the product\_sales. Please note that, there is no orderNumber 16.

```
SELECT * FROM product_sales ORDER BY SalesId;
```

Now will activate the trigger by inserting into orderdetails:

```
INSERT INTO orders  
VALUES(16, '2020-10-01', '2020-10-01', '2020-10-  
01', 'Done', '', 131);  
INSERT INTO orderdetails  
VALUES(16, 'S18_1749', '1', '10', 1);
```

Check product\_sales again, you should have orderNumber 16:

```
SELECT * FROM product_sales ORDER BY SalesId;
```

Note Triggers are not the only way to initiate an ETL process. In fact for performance reasons, it is advised to use the Event engine on large data sets. For more information check: <https://www.mysqltutorial.org/mysql-triggers/working-mysql-scheduled-event/>

## Data marts with Views

With views we can define sections of the datastore and prepare them for a BI operation such as reporting.

View of sales for a specific brand (Vintage\_Cars)

```
DROP VIEW IF EXISTS Vintage_Cars;  
  
CREATE VIEW `Vintage_Cars` AS  
SELECT * FROM product_sales WHERE product_sales.Brand  
= 'Vintage Cars';
```

View of sales in USA:

```
DROP VIEW IF EXISTS USA;  
  
CREATE VIEW `USA` AS  
SELECT * FROM product_sales WHERE country = 'USA';
```

Note the content of Views are generated on-the-fly. For performance reasons, in analytics, so called materialized views are preferred on large data set. This is not supported by MySQL, but there are several ways to implemented. Here is an example: <https://fromdual.com/mysql-materialized-views>

### Exercise4

Create a view, which contains product\_sales rows of 2003 and 2005.